

NetArbiter

-Multi-site Network Emulation -



Hee Won Lee
Cloud Platform Software Research

NetArbiter

INTRODUCTION

NetArbiter Project

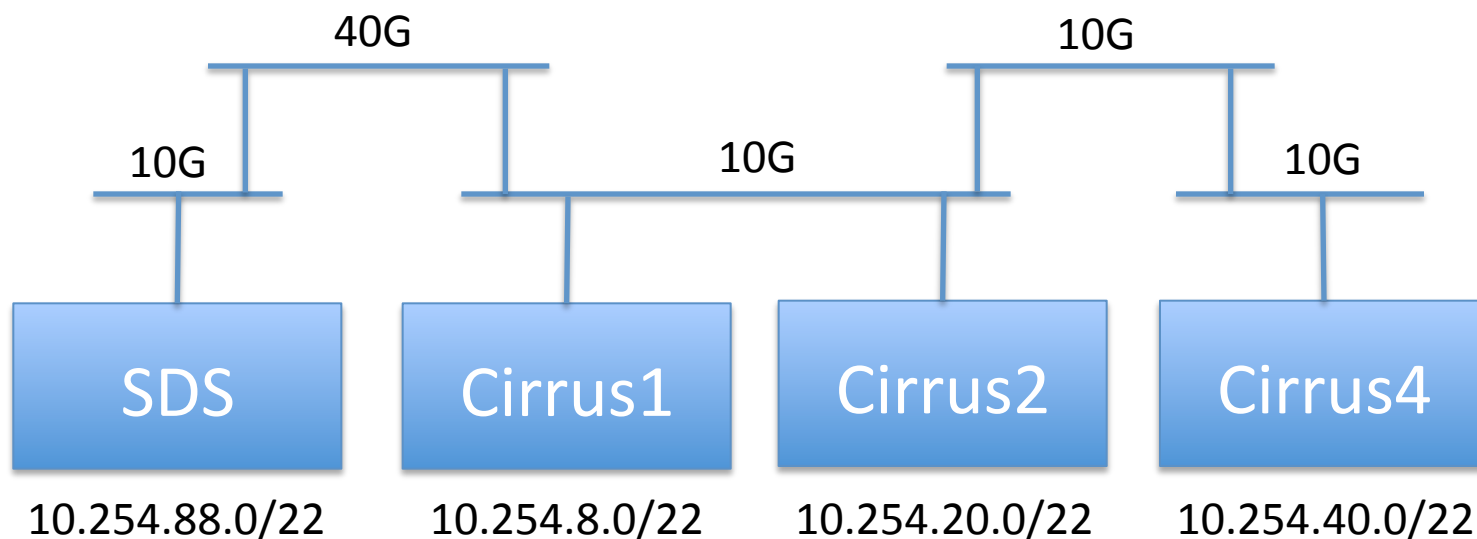
- **Project Description**

- Build a framework for multi-site network emulation

- **Goal**

- Emulate network delay and bandwidth between geographically distributed cloud sites
- Use DevOps for easy implementation and deployment

IP Address Assignment Scheme



- Other clusters
 - nimbus1: 10.254.28.0/22
 - agave: 10.254.0.0/22

Required Components

- Router
 - Will create virtual routers using **OpenStack Neutron**.
- Network emulator
 - Will use **TC** for rate control.
 - Will use **NETEM** for delay emulation.
- Configuration automation
 - Will use a DevOps platform called **Ansible**.

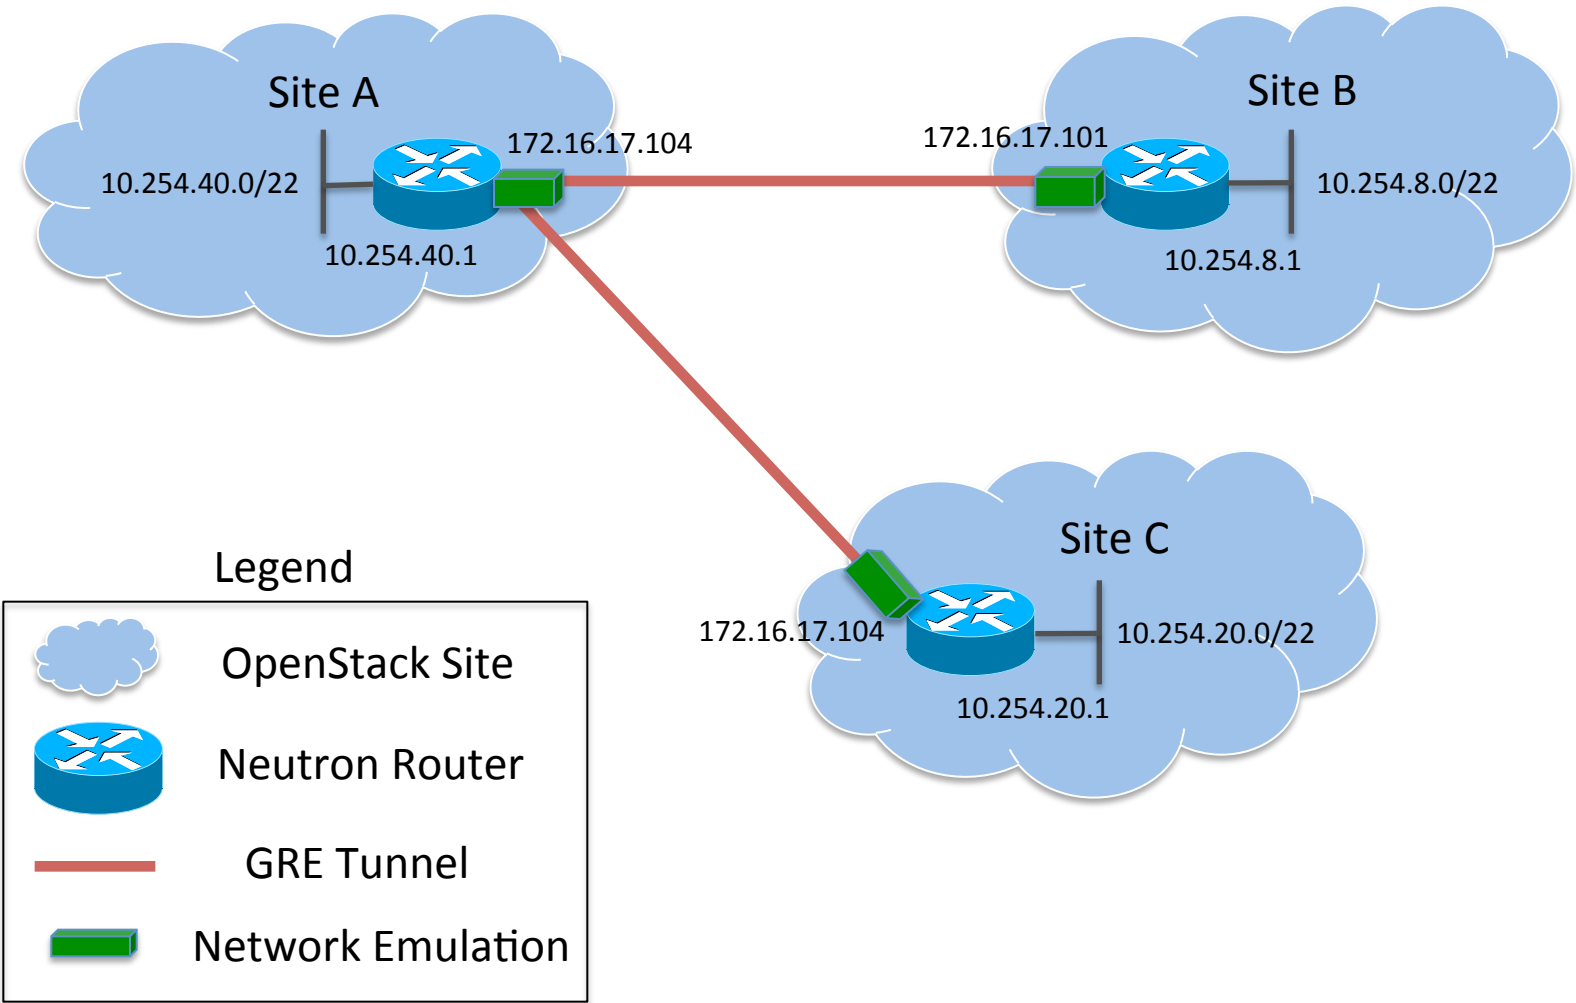
NetArbiter

DESIGN AND IMPLEMENTATION

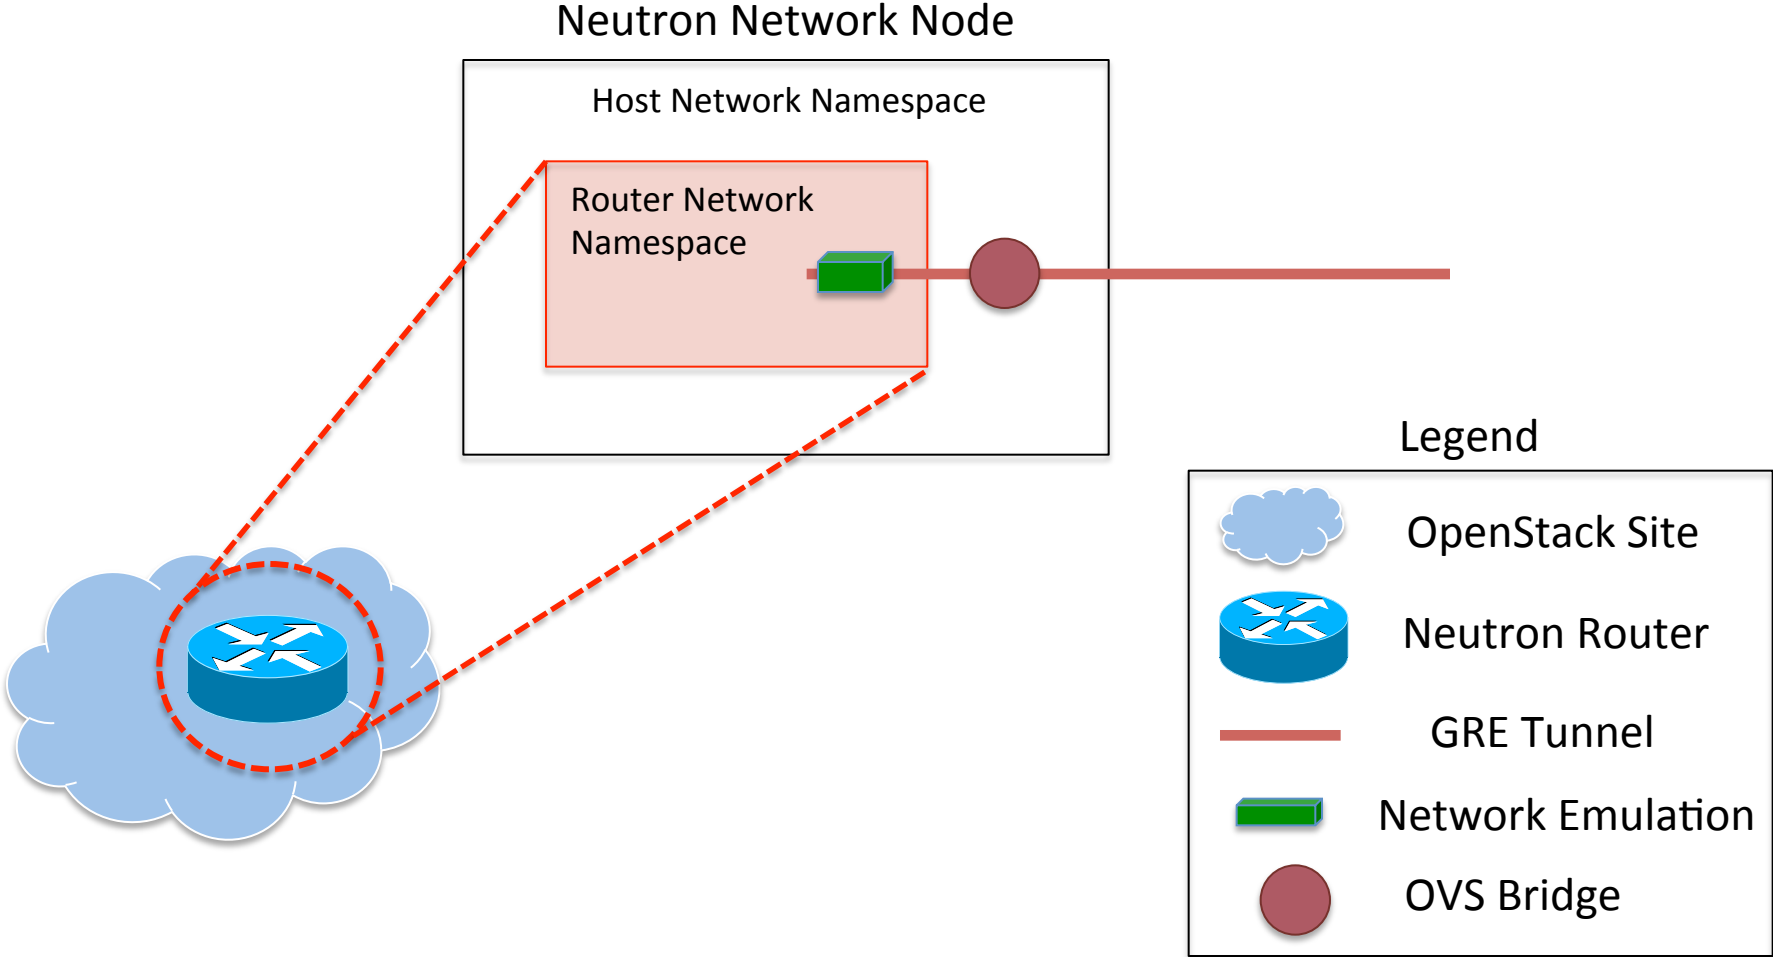
Design & Implementation

- Design
 - Shell scripts
 - Create routers
 - Create tunnels and routes
 - Add network emulation
 - Ansible Playbook
 - Run shell scripts remotely
- Implementation
 - ~ 600 Lines of Code

Configuration Procedure



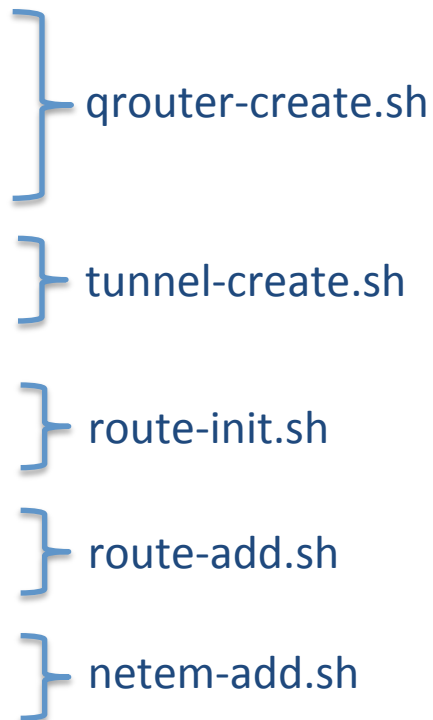
Creating Neutron Router



Linux Network Namespace

- How to list all network namespaces
`# ip netns list`
- How to execute a command in a network namespace
`# ip netns exec $ROUTER_NS $COMMAND`

Shell Scripts

- Create an OVS bridge
 - Create a network, subnet, and router
 - Create a tunnel
 - Set up an interface IP address
 - Add a route
 - Add network emulation
- 
- qrouter-create.sh
- tunnel-create.sh
- route-init.sh
- route-add.sh
- netem-add.sh

qrouter-create.sh

INPUT

```
#!/bin/bash
```

```
OVSVSCTL=/usr/bin/ovs-vsctl
```

```
NEUTRON=/usr/bin/neutron
```

```
$OVSVSCTL add-br $BRIDGE_NAME
```

```
$NEUTRON net-create $NET_NAME --shared
```

```
$NEUTRON subnet-create $NET_NAME $SITE_CIDR --name $SUBNET_NAME --gateway $SITE_GATEWAY
```

```
$NEUTRON router-create $ROUTER_NAME
```

```
$NEUTRON router-interface-add $ROUTER_NAME $SUBNET_NAME
```

- BRIDGE_NAME
- NET_NAME
- SUBNET_NAME, SITE_CIDR, SITE_GATEWAY
- ROUTER_NAME

tunnel-create.sh

INPUT

- BRIDGE_NAME
- TUNNEL_NAME
- REMOTE_IP

```
#!/bin/bash
```

```
OVSVSCTL=/usr/bin/ovs-vsctl
```

```
$OVSVSCTL add-port $BRIDGE_NAME $TUNNEL_NAME -- set interface $TUNNEL_NAME type=gre  
options:remote_ip=$REMOTE_IP
```

route-init.sh

```
#!/bin/bash
```

```
IP=/sbin/ip
```

```
IFCONFIG=/sbin/ifconfig
```

```
$IP link set $BRIDGE_NAME netns $ROUTER_NS
```

```
$IP netns exec $ROUTER_NS $IFCONFIG $BRIDGE_NAME $WANIP_PREFIXLEN
```

INPUT

- BRIDGE_NAME
- ROUTER_NS
- WANIP_PREFIXLEN

route-add.sh

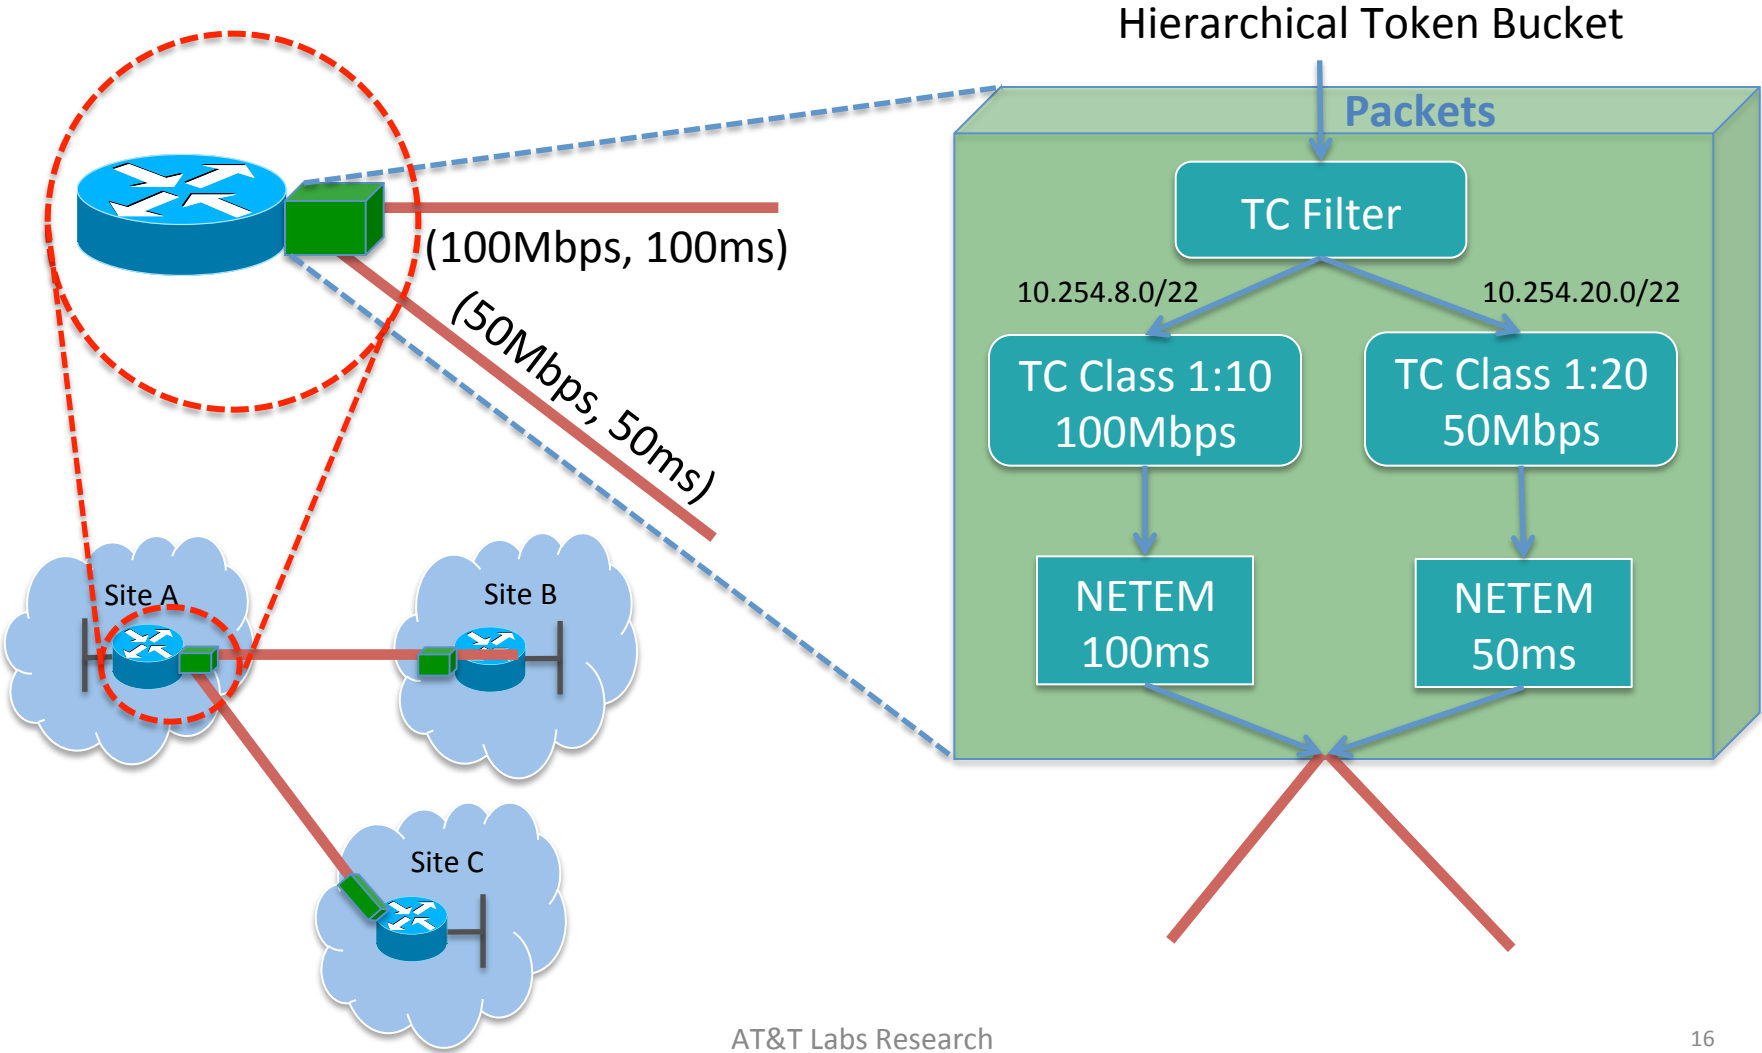
INPUT

```
#!/bin/bash  
IP=/sbin/ip  
ROUTE=/sbin/route
```

```
- ROUTER_NS  
- REMOTE_CIDR  
- REMOTE_GATEWAY
```

```
$IP netns exec $ROUTER_NS $ROUTE add -net $REMOTE_CIDR gw $REMOTE_GATEWAY
```

Network Emulation



netem-add.sh

```
#!/bin/bash
IP=/sbin/ip
TC=/sbin/tc
NEUTRON=/usr/bin/neutron
```

INPUT

- ROUTER_NS
- BRIDGE_NAME
- **RATE, CEIL**
- **DELAY**
- **REMOTE_CIDR**

```
$IP netns exec $ROUTER_NS $TC qdisc replace dev $BRIDGE_NAME handle 1: root htb default 10
```

```
$IP netns exec $ROUTER_NS $TC class replace dev $BRIDGE_NAME parent 1: classid
1:$TC_CLASSID htb rate $RATE ceil $CEIL
```

```
$IP netns exec $ROUTER_NS $TC qdisc replace dev $BRIDGE_NAME parent 1:$TC_CLASSID handle
$TC_CLASSID: netem delay $DELAY
```

```
$IP netns exec $ROUTER_NS $TC filter replace dev $BRIDGE_NAME parent 1: protocol ip prio 1
u32 match ip dst $REMOTE_CIDR flowid 1:$TC_CLASSID
```

To Clean Up

- Scripts for clean-ups
 - qrouter-destroy.sh
 - tunnel-destroy.sh
 - route-del.sh
 - netem-del.sh

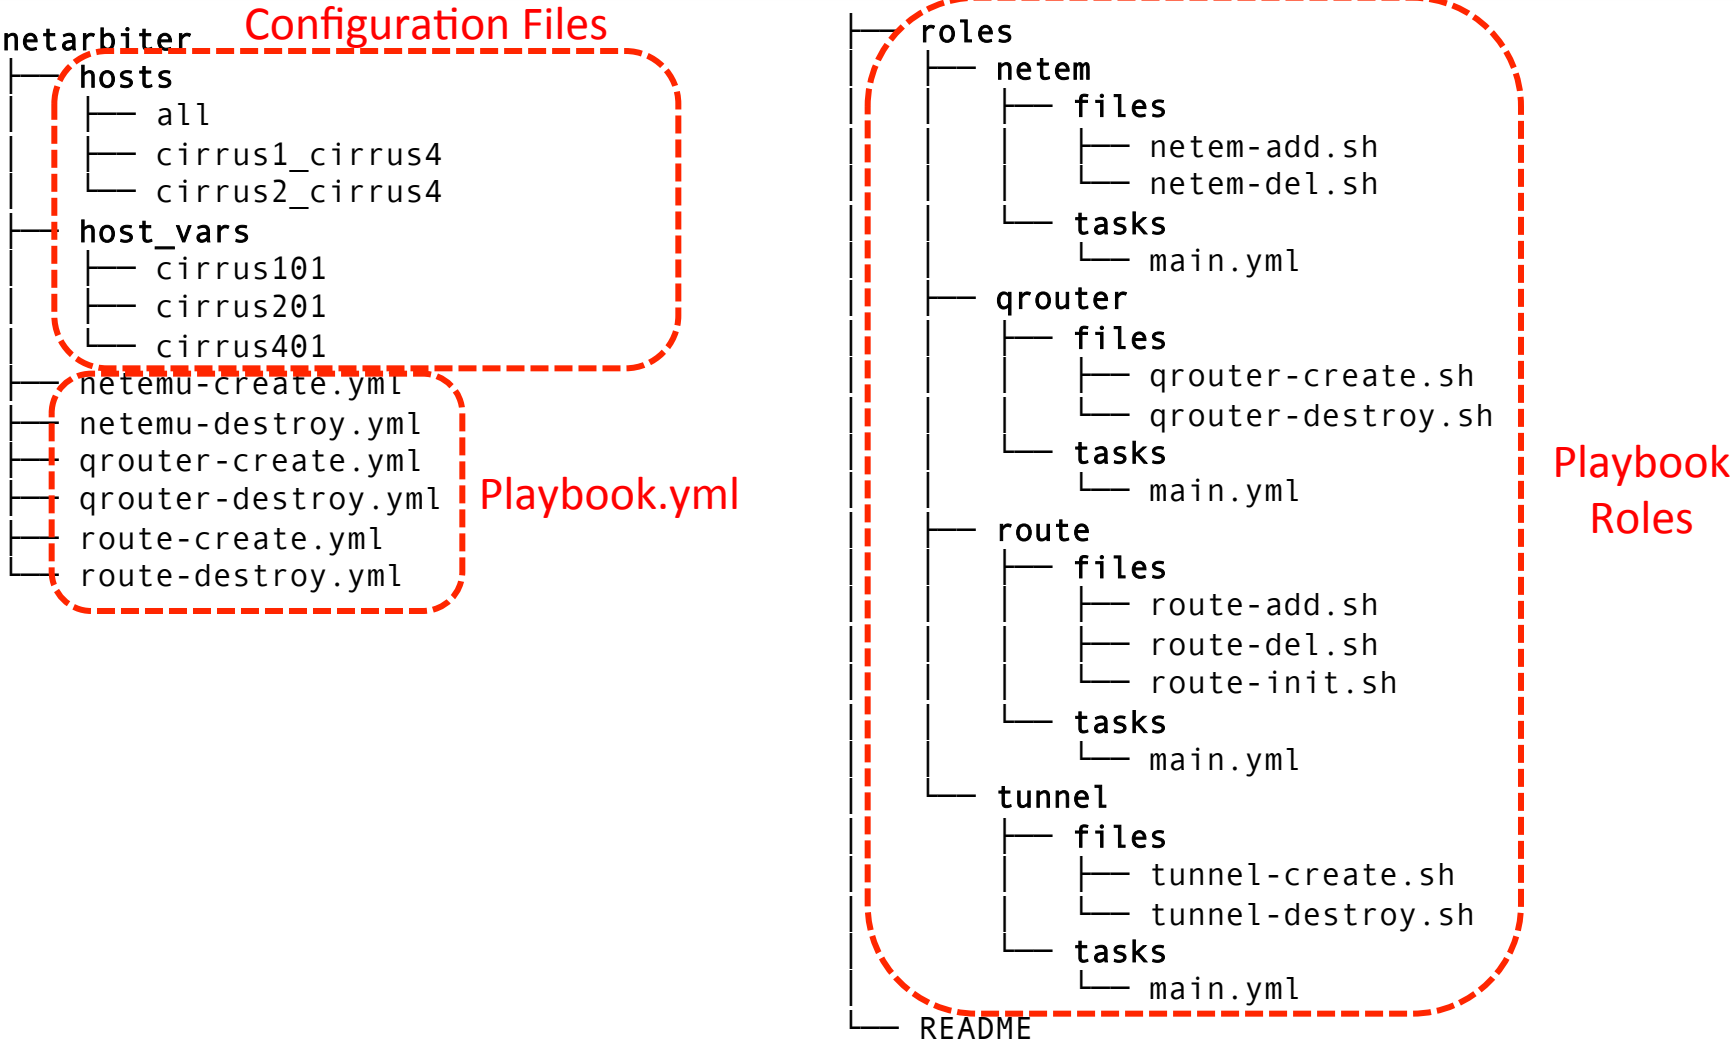
NetArbiter

INTEGRATION WITH ANSIBLE

What is Ansible?

- What is Ansible?
 - *Ansible* is an IT automation engine that automates cloud provisioning, configuration management, application deployment, intra-service orchestration, etc.
 - Red Hat announced the acquisition of *Ansible* on October 16, 2015.
- Features
 - Uses no agents and no additional custom security infrastructure.
 - Easy deployment
 - Uses YAML in the form of Ansible Playbooks
 - Automation jobs in YAML

NetAribiter Directory Structure



Link Configuration (ex. cirrus1_cirrus4)

```
[multisite_link]
```

```
cirrus101 REMOTE_IP=135.197.240.47 REMOTE_CIDR=10.254.40.0/22
```

```
REMOTE_GATEWAY=172.16.17.104
```

```
cirrus401 REMOTE_IP=135.197.240.11 REMOTE_CIDR=10.254.8.0/22
```

```
REMOTE_GATEWAY=172.16.17.101
```

```
[multisite_link:vars]
```

```
DELAY=80ms
```

```
RATE=10mbit
```

```
CEIL=100mbit
```

```
TUNNEL_NAME=gre_cirrus1_cirrus4
```



Host Configuration (ex. cirrus101)

WANIP_PREFIXLEN: 172.16.17.101/24

SITE_CIDR: 10.254.8.0/22

SITE_GATEWAY: 10.254.8.1

BRIDGE_NAME: br-multisite

NET_NAME: multisite-net

SUBNET_NAME: multisite-subnet

ROUTER_NAME: multisite-router

Neutron-related

os_env:

OS_PROJECT_NAME: admin

OS_TENANT_NAME: admin

OS_USERNAME: admin

OS_PASSWORD: *****

OS_AUTH_URL: http://controller:35357

OS_VOLUME_API_VERSION: 2

OS_IMAGE_API_VERSION: 2



Playbook: qrouter-create.yml

```
---  
- hosts: "{{ HOST }}"  
  remote_user: yourusername  
  sudo: yes  
  gather_facts: no  
  
  roles:  
    - {role: qrouter, subcommand: qrouter_create}
```


Playbook: route-create.yml

- hosts: multisite_link
 - remote_user: yourusername
 - sudo: yes
 - gather_facts: no
- roles:
 - {role: tunnel, subcommand: tunnel_create}
 - {role: route, subcommand: route_init}
 - {role: route, subcommand: route_add}

Playbook: netem-create.yml

```
- hosts: multisite_link
  remote_user: yourusername
  sudo: yes
  gather_facts: no

roles:
  - {role: netem, subcommand: netem_add}
```

Playbook Role: roles/netem/main.yml

```
---
- name: Add network emulation - netem-add.sh
  script: netem-add.sh {{ BRIDGE_NAME }} {{ ROUTER_NAME }} {{ REMOTE_CIDR }}
  {{ DELAY }} {{ RATE }} {{ CEIL }}
  environment: os_env
  when: subcommand == 'netem_add'
  register: netem_add

- name: Delete network emulation - netem-del.sh
  script: netem-del.sh {{ BRIDGE_NAME }} {{ ROUTER_NAME }} {{ REMOTE_CIDR }}
  environment: os_env
  when: subcommand == 'netem_del'
  register: netem_del

- debug: var=netem_add.stdout_lines
- debug: var=netem_del.stdout_lines
```

How to Run

0. Setup

- Store all hosts' domain names (or IP addresses) to an inventory file "all".
- Create an inventory file for each host.
- Create an inventory file for each link between two hosts.

1. Create/destroy a qrouter (and an ovs bridge) in a host

```
$ ansible-playbook -K -i hosts/all qrouter-create.yml -e HOST=<hostname>  
$ ansible-playbook -K -i hosts/all qrouter-destroy.yml -e HOST=<hostname>
```

2. Create/destroy routes (and tunnels) to connect two sites

```
$ ansible-playbook -K -i hosts/<inventory> route-create.yml  
$ ansible-playbook -K -i hosts/<inventory> route-destroy.yml
```

3. Create delay, rate, and ceil with TC/NETEM

```
$ ansible-playbook -K -i hosts/<inventory> netemu-create.yml [-e DELAY=<delay> -e  
RATE=<rate> CEIL=<ceil>]  
$ ansible-playbook -K -i hosts/<inventory> netemu-destroy.yml [-e DELAY=<delay> -e  
RATE=<rate> CEIL=<ceil>]
```

How to Run

1. Create virtual routers

```
$ ansible-playbook -K -i hosts/all qrouter-create.yml -e HOST=cirrus401
```

```
$ ansible-playbook -K -i hosts/all qrouter-create.yml -e HOST=cirrus201
```

→ Check the Network tab of OpenStack Horizon to see if routers are created.

2. Create routes between virtual routers

```
$ ansible-playbook -K -i hosts/cirrus2_cirrus4 route-create.yml
```

3. Launch two instances from cirrus2 and cirrus4 via OpenStack Horizon

→ Check connectivity by running ping from instances.

4. Add delay

```
$ ansible-playbook -K -i hosts/cirrus2_cirrus4 netemu-create.yml
```

```
$ ansible-playbook -K -i hosts/cirrus2_cirrus4 netemu-create.yml -e DELAY=80ms
```

→ Check delays by running 'ping 10.254.20.3' from a cirrus4's instance (i.e. 10.254.40.3).



How to Clean Up

Clean up all configurations

```
$ ansible-playbook -K -i hosts/all qrouter-destroy.yml -e HOST=cirrus201
```

```
$ ansible-playbook -K -i hosts/all qrouter-destroy.yml -e HOST=cirrus401
```

Note

- In order to destroy Neutron routers, first terminal all instances connected to the routers, or detach all connections to the routers.
- *qrouter-destroy.sh* deletes router, subnet, net from Neutron and bridge from OVS.

qrouter-destroy.sh

```
#!/bin/bash
OVSVSCTL=/usr/bin/ovs-vsctl
NEUTRON=/usr/bin/neutron

$NEUTRON router-interface-delete $ROUTER_NAME $SUBNET_NAME
$NEUTRON router-delete $ROUTER_NAME
$NEUTRON subnet-delete $SUBNET_NAME
$NEUTRON net-delete $NET_NAME
$OVSVSCTL del-br $BRIDGE_NAME
```

How to Debug

0. Connect via ssh to the Neutron network node of a cloud site

1. Check the creation of an OVS bridge

```
# ovs-vsctl show
```

2. Find *router_id* of a newly-created virtual router

```
# neutron router-list
```

3. Find the network namespace of a virtual router (i.e., *qrouter-router_id*)

```
# ip netns list
```

4. Check the routing table of a virtual router

```
# ip netns exec qrouter-router_id netstat -rn
```

4. Check TC/NETEM configuration

```
# ip netns exec qrouter-router_id tc qdisc show
```

```
# ip netns exec qrouter-router_id tc class show dev bridge_name
```

```
# ip netns exec qrouter-router_id tc filter show dev bridge_name
```